

Supercharge your Python library using ast parsing



[Presentation Materials](#)

PyCon US 2025

Pittsburgh, PA

Adam Glustein



Code from [Point72/csp](#)

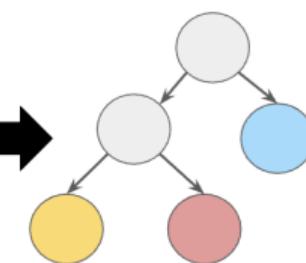
Main idea

- Performant Python libraries are usually implemented in **compiled** languages (C/C++/Rust)
- Between the **user's** Python code and the **library's** underlying implementation, we can transform the ast to squeeze out even more performance in the "hot path"

Main idea

```
import mylibrary

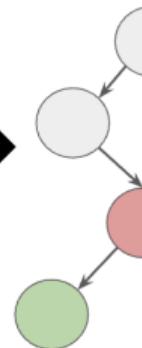
@mylibrary_decorator
def user_function(x):
    ...do something...
    mylibrary.func(x)
    return something
```



User Python code

```
#include <mylibrary.h>

static
PyObject * impl()
{
    ...do something...
    return result;
}
```



Modified AST

Compiled implementation

These are *practical* optimizations

- Strategies presented are taken from **csp**, an open-source Python stream processing library by Point72
- Links to the code are provided in case you don't believe me!



Agenda

Background

- ① What is an abstract syntax tree and how can we modify the *ast*?
- ② Brief overview of computation graphs and CSP

The fun stuff

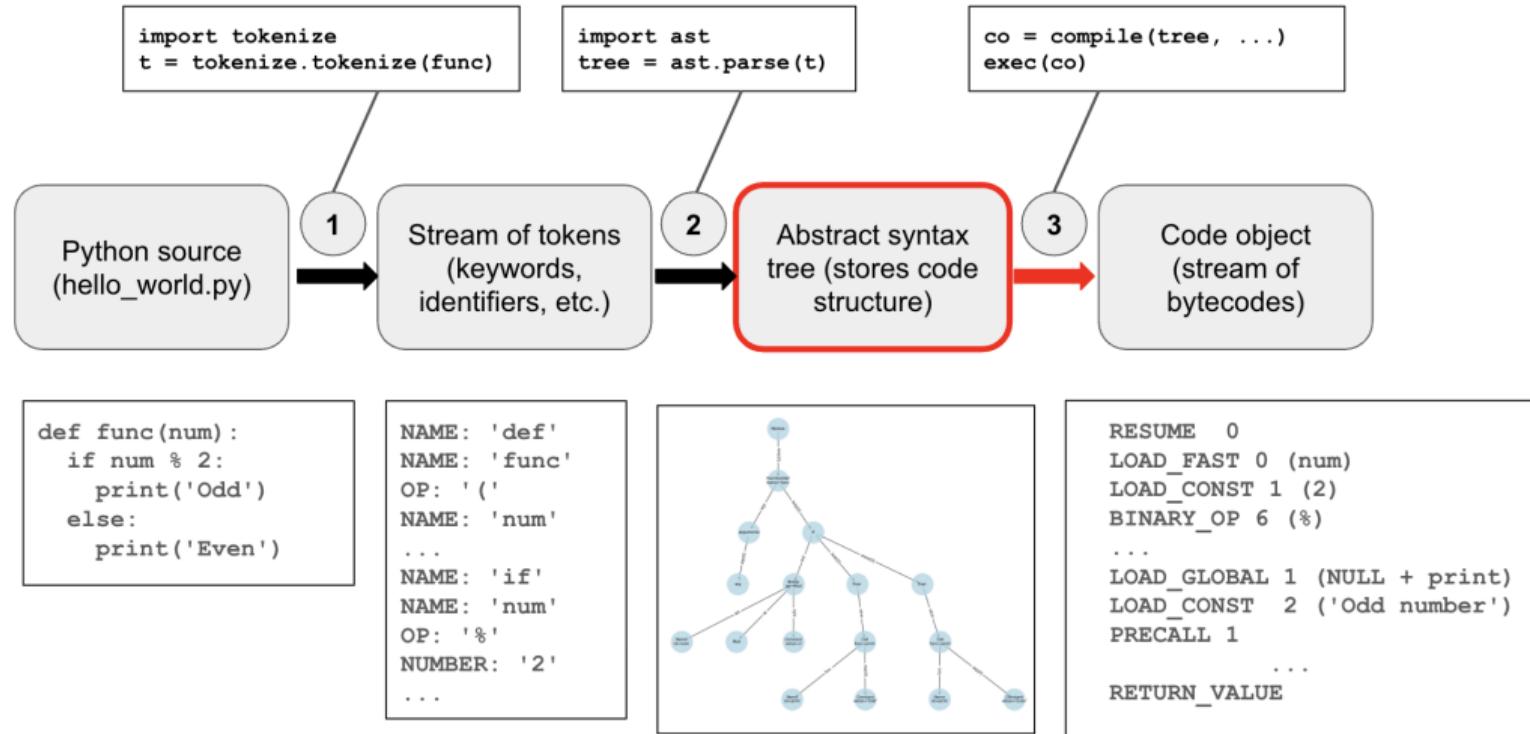
- ① Transforming functions to generators in the AST
- ② Reducing member function call overhead with AST transformation
- ③ Fast local access of C/C++ data using the AST

A bit about myself

- Quant Developer at **Point72** in New York
- I help maintain the csp project
- We do a lot of cool stuff with the *ast*!



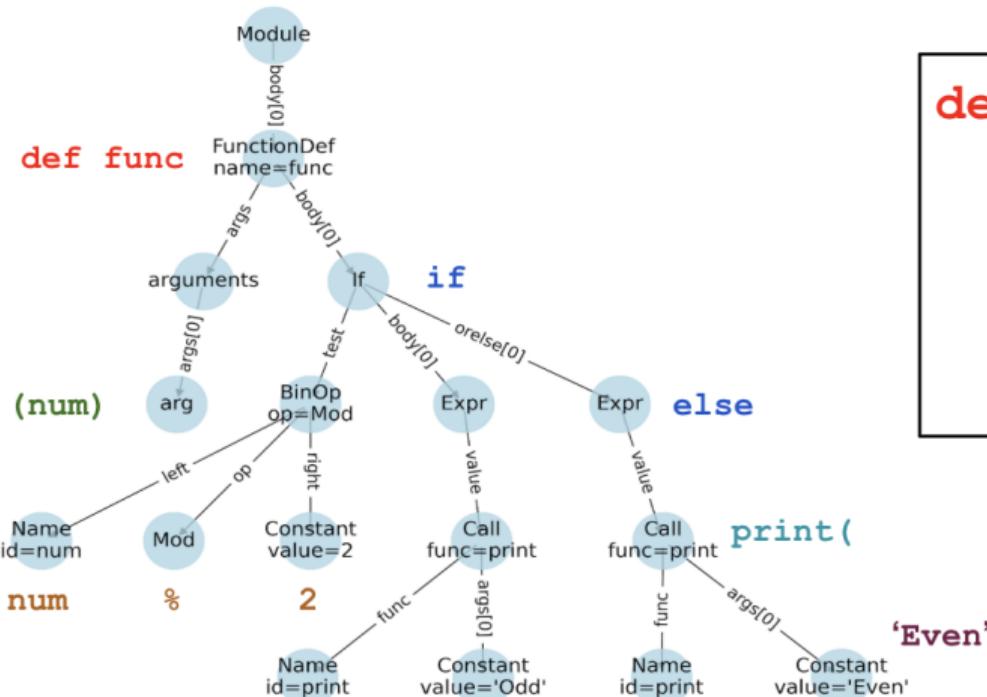
How Python code is executed



The Abstract Syntax Tree

- Represents the **structure** of a program
- *Abstracts* away unimportant information from the raw tokens (parentheses, newlines etc.)
- Last intermediate representation before bytecode generation
- Create the tree object in Python using `ast.parse`

The Abstract Syntax Tree



```
def func(num):  
    if num % 2:  
        print('Odd')  
    else:  
        print('Even')
```

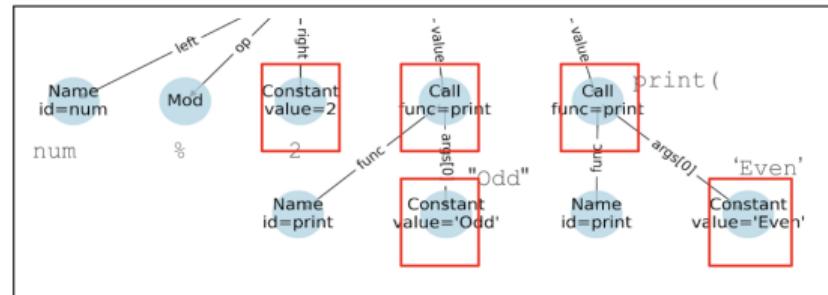
Let's visit an ast

- Subclass `ast.NodeVisitor` to traverse an AST

```
1 import ast
2
3 class MyVisitor(ast.NodeVisitor):
4     def visit_Call(self, node):
5         if isinstance(node.func, ast.Name):
6             print(f"Function call: {node.func.id}")
7             self.generic_visit(node) # continue traversal
8
9     def visit_Constant(self, node):
10        print(f"Constant with value: {node.value}")
11        self.generic_visit(node) # continue traversal
```

Let's visit an ast

```
1 import inspect  
2  
3 def func(num):  
4     if num % 2:  
5         print("Odd")  
6     else:  
7         print("Even")  
8  
9 src = inspect.getsource(func)  
10 tree = ast.parse(src)  
11 MyVisitor().visit(tree)
```



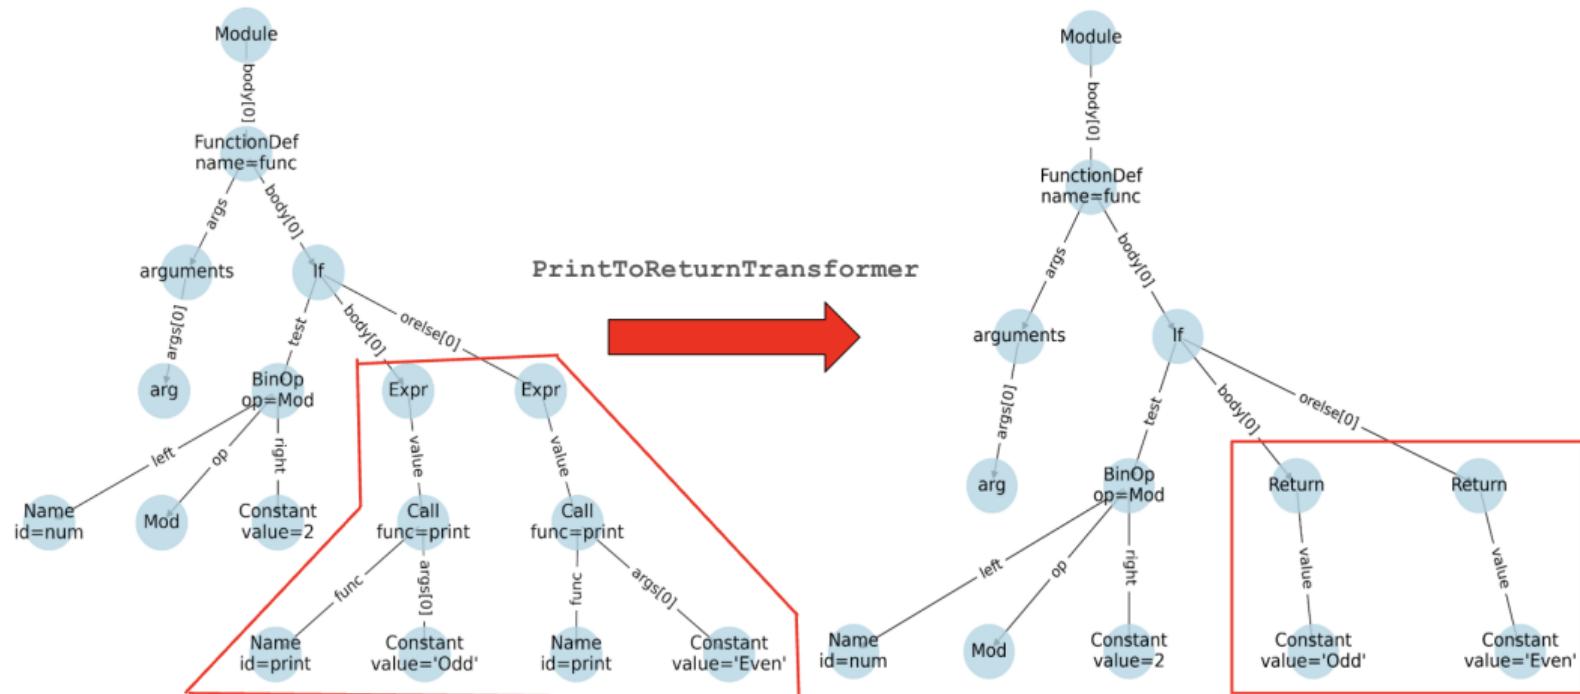
Constant with value: 2
Function call: print
Constant with value: Odd
Function call: print
Constant with value: Even

Let's transform an ast

- Subclass `ast.NodeTransformer` to transform an AST

```
1 import ast
2
3 class PrintToReturnTransformer(ast.NodeTransformer):
4     def visit_Expr(self, node):
5         # Check if this expr contains a print(...)
6         if isinstance(node.value, ast.Call)
7             and node.value.func.id == "print":
8                 # replace print with return
9                 return ast.Return(value=node.value.args[0])
10            # Continue normally for other functions
11            return self.generic_visit(node)
```

Let's transform an ast



Let's transform an ast

```
1 tree = ast.parse(func_code)
2 new_tree = PrintToReturnTransformer().visit(tree)
3
4 # Fix line numbers and other details
5 ast.fix_missing_locations(new_tree)
6
7 # Compile and execute
8 namespace = {}
9 co = compile(new_tree, "<ast>", mode="exec")
10 exec(co, namespace)
11
12 # Test the new function
13 func = namespace["func"]
14 func(5) # returns "Odd"
```

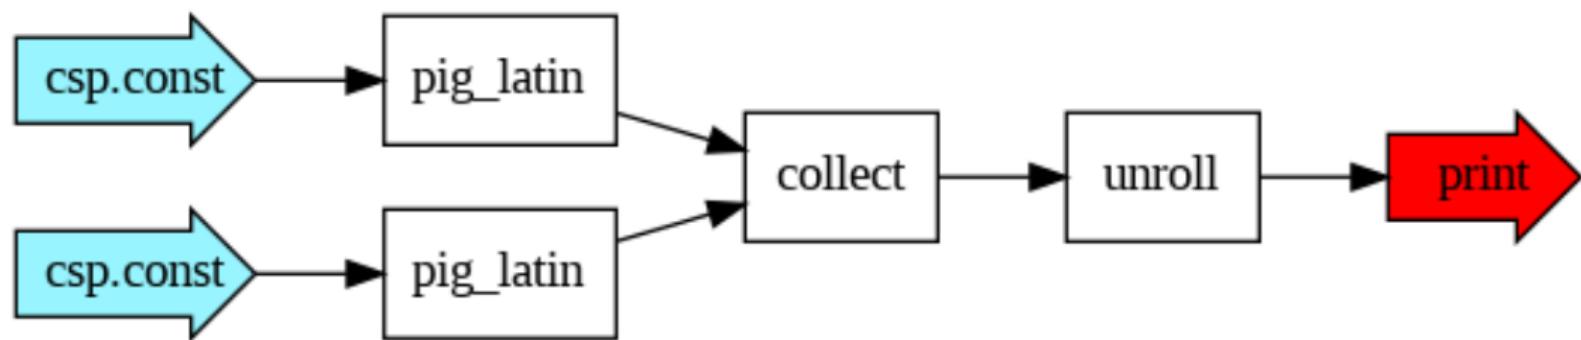
Key points to remember about the AST

- AST is an **abstract** representation of code
- Intermediate stage between Python source code and bytecode instructions
- We can use standard `ast` library to **transform** an AST

Next, we're going to look at something that *seems* entirely unrelated...

Computation graphs

- Directed, acyclic graph (DAG) that executes some task
- Functions → nodes, inputs → sources, outputs → sinks
- Used in many common Python libraries (torch, tensorflow, dask, airflow, etc.)



Computation graph that transforms words to pig latin

A (very) brief overview of CSP

- CSP = Composable Stream Processing
- Computation graph for **real-time streaming** data
 - Financial markets, climate analysis, transportation/supply-chain management, etc.
- Users create **nodes**, compose them into **graphs**, then compose those into larger graphs
- Engine code written in **C++** and exposed via C API

How CSP is implemented

- Users write nodes in Python that manipulate **time series** data
- Nodes are invoked when their inputs "**tick**"
- Nodes call `csp.output` to "**tick out**" a value
- `csp.node` and time series (`csp.ts`) correspond to objects in C++
- Runtime lives in C++, including the values of all `csp.ts` objects

@csp.node

- Functions that execute as part of a graph
- Inputs and outputs are **time-series** data (`csp.ts[T]`)
- **Stateful** and **atomic** units of computation

```
1 import csp
2
3 @csp.node
4 def sum(input: csp.ts[int]) -> csp.ts[int]:
5     with csp.state():
6         s_sum = 0 # initial state
7
8     if csp.ticked(input): # "ticked" = x has a new value
9         s_sum += input
10        csp.output(s_sum) # "tick" out an output
```

A closer look at the sum node

- Under the hood, there are four different AST transformations we'll look at

```
@csp.node

def sum(input: csp.ts[int]) -> csp.ts[int]:
    with csp.state():
        s_sum = 0 # initial state
    if csp.ticked(input):
        s_sum += input
    csp.output(s_sum)
```

- 1 Transformed to generator
- 2 Member function substitution
- 3 Member function substitution
- 4 Fast local variable access

1) Making nodes stateful functions

```
with csp.state():
    s_sum = 0 # initial state
```

- Nodes hold **state** between cycles
- **Generators** are useful constructs for maintaining state between cycles
- Transform nodes to generators using ast transformation:
 - Wrap function body with `ast.While`
 - `ast.Return` → `ast.Yield`

2) Reducing member function call overhead: no arguments

```
if csp.ticked(input): # input.ticked()
```

- `csp.ticked`: whether an input has a new value this cycle
- AST logic translates `ticked` to a **member function** on the input (i.e. `input.ticked()`) that is defined in C++
- `csp.ticked` is called *very often*, so how can we optimize it?

2) Reducing member function call overhead: no arguments

- PyTypeObject in C/C++ defines commonly called functions (tp_init, tp_repr etc.)
- We can substitute the ticked call for a type method that avoids lookup of named attribute
- AST transformation in csp allows us to hide this implementation from users

obj.ticked()

```
0 LOAD_FAST          0 (obj)
2 LOAD_ATTR          0 (ticked)
4 CALL_FUNCTION      0
6 RETURN_VALUE
```

+obj

```
0 LOAD_FAST          0 (obj)
2 UNARY_POSITIVE
4 RETURN_VALUE
```

3) Reducing member function call overhead: with arguments

```
csp.output(s_sum) # <out_ts>.output(s_sum)
```

- Same situation as before, except we have an argument
- We transform `output` to a binary add (`ast.BinOp`)

`obj.output(1)`

```
LOAD_FAST      0  (obj)
LOAD_ATTR      1  (output + NULL|self)
LOAD_CONST     1  (1)
CALL           1
RETURN_VALUE
```

`obj+1`

```
LOAD_FAST      0  (obj)
LOAD_CONST     1  (1)
BINARY_OP      0  (+)
RETURN_VALUE
```

3) Reducing member function call overhead: with arguments

- Simple microbenchmark to show performance improvement

```
static PyObject* MyClass_output(PyObject* self, PyObject* other) {
    Py_INCREF(other);
    return other;
}

// add to method table as normal member function
static PyMethodDef MyClass_methods[] = {
    {"output", (PyCFunction) MyClass_output, METH_VARARGS, ""},
    ...
};

// also define the same method as BINARY_ADD
static PyNumberMethods MyClass_as_number = {
    (binaryfunc) MyClass_output,
    ...
};
```

3) Reducing member function call overhead: with arguments

```
IPython 8.12.2 -- An enhanced Interactive Python. Type '?' for help.

In [1]: from mymodule import MyClass

In [2]: obj = MyClass()

In [3]: %timeit obj+1
18 ns ± 1.36 ns per loop (mean ± std. dev. of 7 runs, 100,000,000 loops each)

In [4]: %timeit obj.output(1)
34.5 ns ± 0.176 ns per loop (mean ± std. dev. of 7 runs, 10,000,000 loops each)
```

Python Version	obj.output(1)	obj+1	Improvement
3.8.12	34.5 ns	18 ns	1.92x
3.11.12	30.6 ns	16.4 ns	1.87x
3.13.2	30.4 ns	19.8 ns	1.54x

4) Fast local variable access of C/C++ data

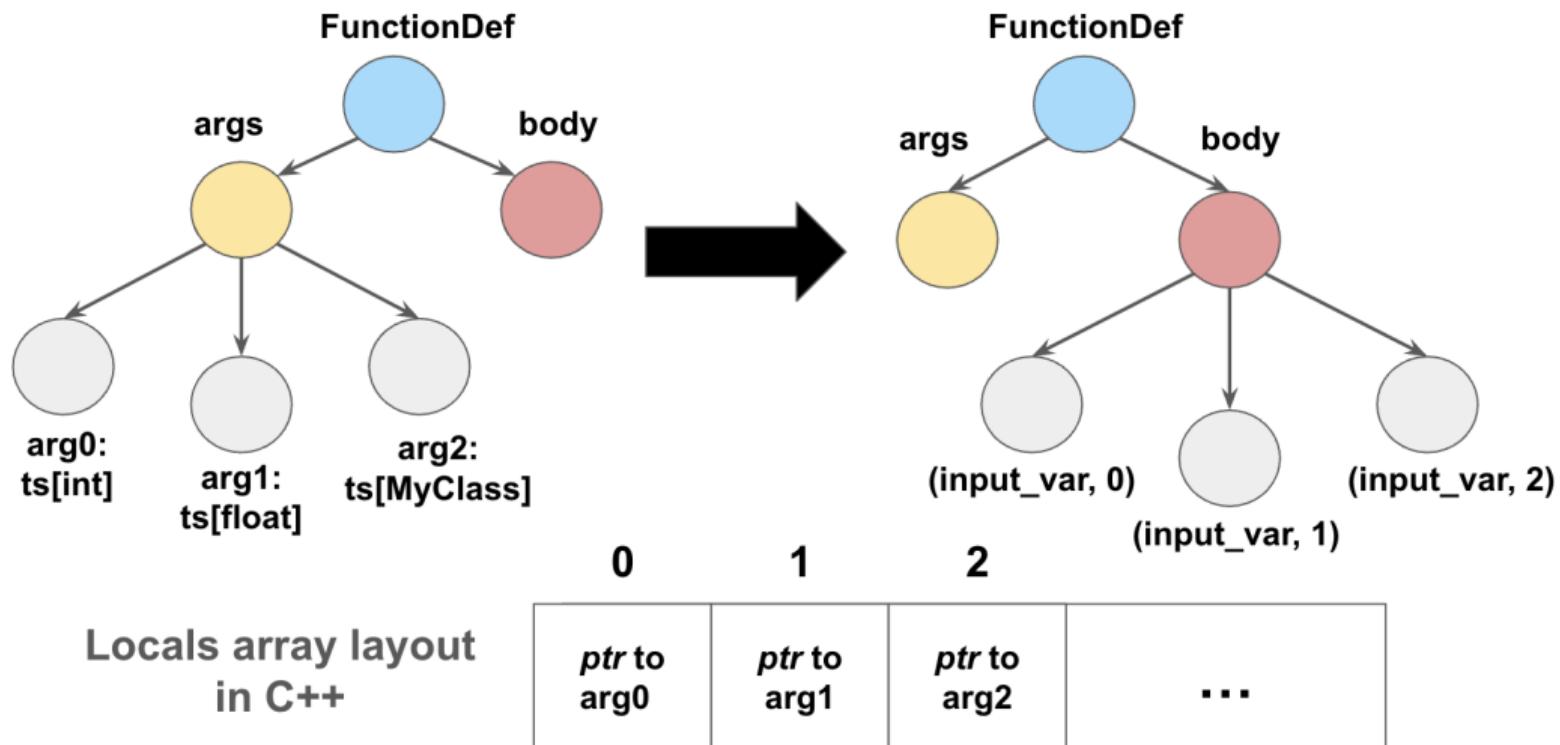
```
s_sum += input
```

- Value of all `csp.ts` objects live in C++ (i.e. `csp.ts[int]` stored as `int64_t` member on `input` object)
- Users need to access those values in their Python nodes
- This needs to be *as fast as possible!*

4) Fast local variable access of C/C++ data

- Take a variable defined in C++ and store it directly in the **fast locals array**
- We grab pointers to the input variables on the stack and **store those in C++**
- Use the ast to index each node's inputs, then store the variable with its **stack index in an ast.Tuple**
- `input → ("input_var", idx)`

4) Fast local variable access of C/C++ data

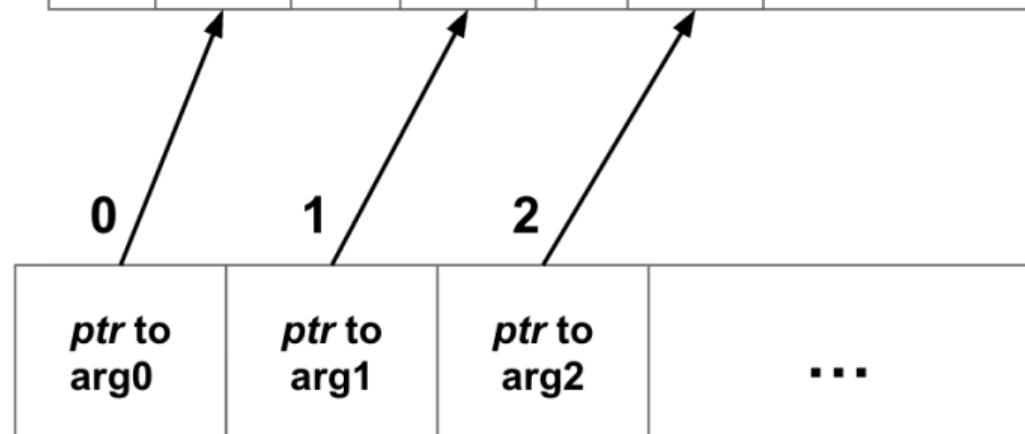


4) Fast local variable access of C/C++ data

Locals array in
Python stack frame



Locals array layout
in C++



4) Fast local variable access of inputs

- Then, when we invoke the node at **runtime** we get LOAD_FAST on each csp.ts

```
46 CALL_INTRINSIC_1          5 (INTRINSIC_UNARY_POSITIVE)
48 POP_JUMP_IF_FALSE        10 (to 70)

15   50 LOAD_FAST             4 (s_sum)
      52 LOAD_FAST             3 (input)
      54 BINARY_OP             13 (+=)
      58 STORE_FAST            4 (s_sum)

16   60 LOAD_FAST             2 (#outp_0)
      62 LOAD_FAST             4 (s_sum)
      64 BINARY_OP             0 (+)
```

The final product

```
@csp.node
def sum(input: csp.ts[int])
    -> csp.ts[int]:
        with csp.state():
            s_sum = 0 # initial state

        if csp.ticked(input):
            s_sum += input
            csp.output(s_sum)
```

Modify
AST



```
def sum():
    #inp_0 = 'input_proxy', 0
    #outp_0 = 'output_proxy', 0
    input = 'input_var', 0
    yield
    del input
    s_sum = 0
    while True:
        yield
        if +#inp_0:
            s_sum += input
            #outp_0 + s_sum
```

Conclusion

- Transforming the ast is a powerful (but seldom used) strategy to boost the performance of Python libraries that interface with compiled code
- Modify a user's Python code to best fit the C/C++ implementation without compromising the API
- Specific techniques shown here (from `csp`) run in critical, high-performance production code